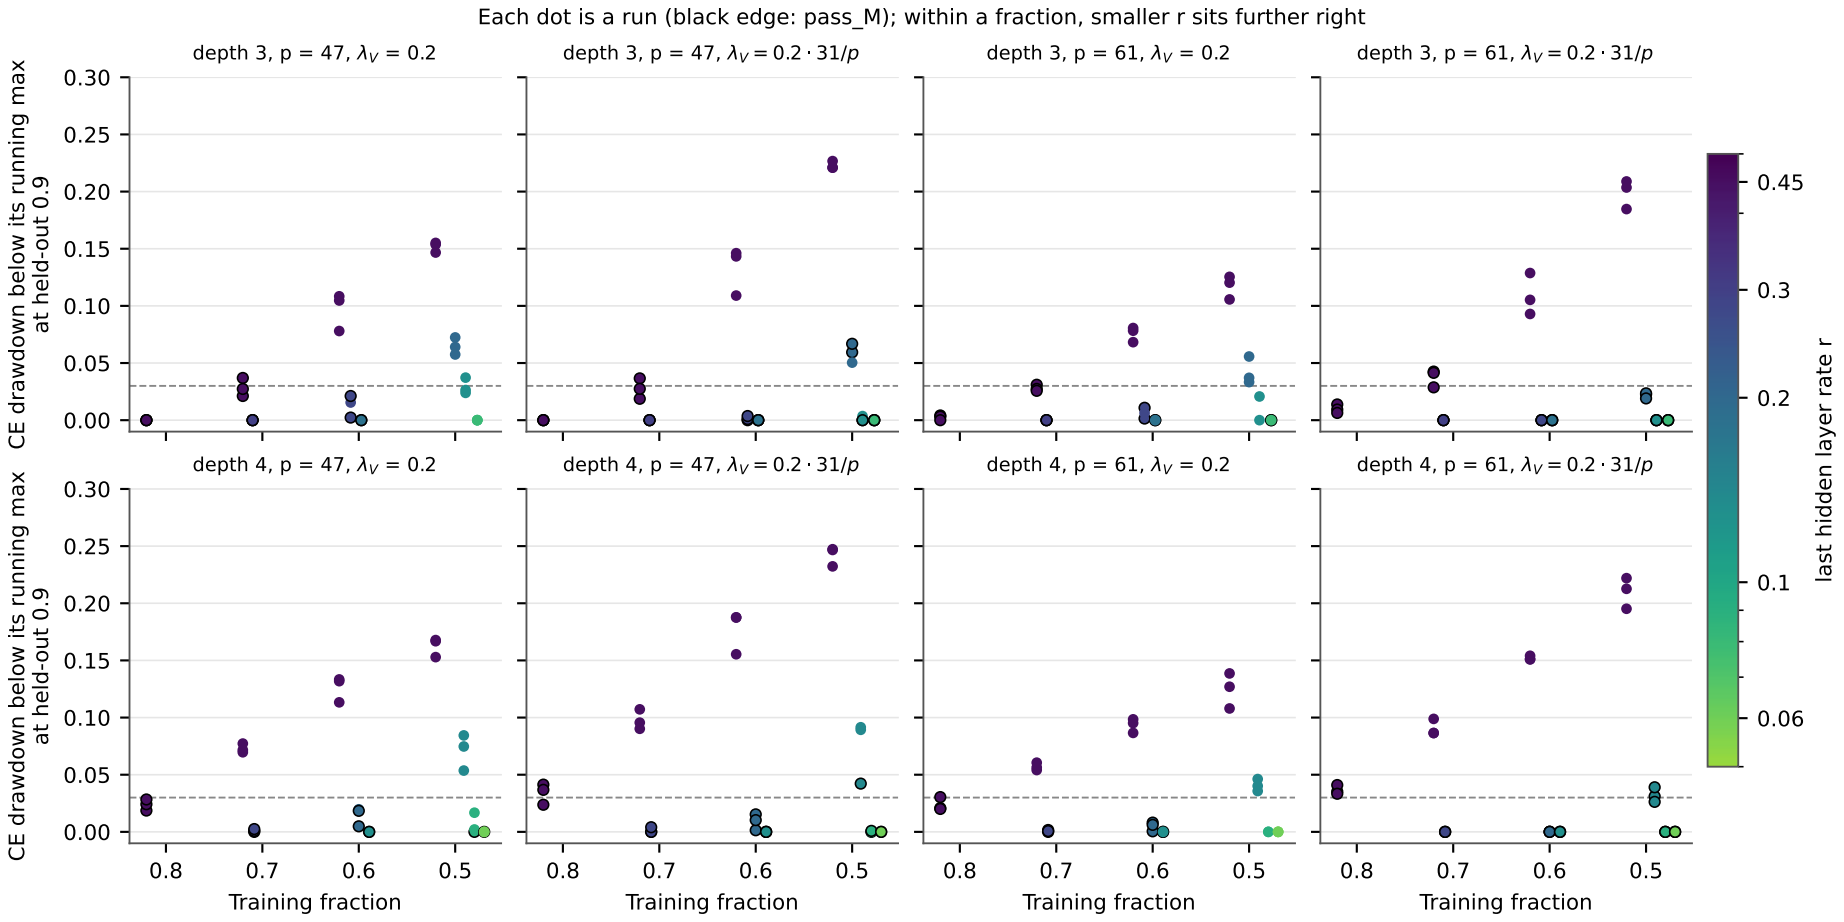
Each dot is a run (black edge: pass_M); within a fraction, smaller r sits further right
depth 3, p = 47, $\lambda_V = 0.2$
depth 3, p = 47, $\lambda_V = 0.2 \cdot 31/p$
depth 3, p = 61, $\lambda_V = 0.2$
depth 3, p = 61, $\lambda_V = 0.2 \cdot 31/p$
depth 4, p = 47, $\lambda_V = 0.2$
depth 4, p = 47, $\lambda_V = 0.2 \cdot 31/p$
depth 4, p = 61, $\lambda_V = 0.2$
depth 4, p = 61, $\lambda_V = 0.2 \cdot 31/p$
CE drawdown below its running max at held-out 0.9
0.30
0.25
0.20
0.15
0.10
0.05
0.00
0.8
0.7
0.6
0.5
Training fraction
last hidden layer rate r
0.45
0.3
0.2
0.1
0.06