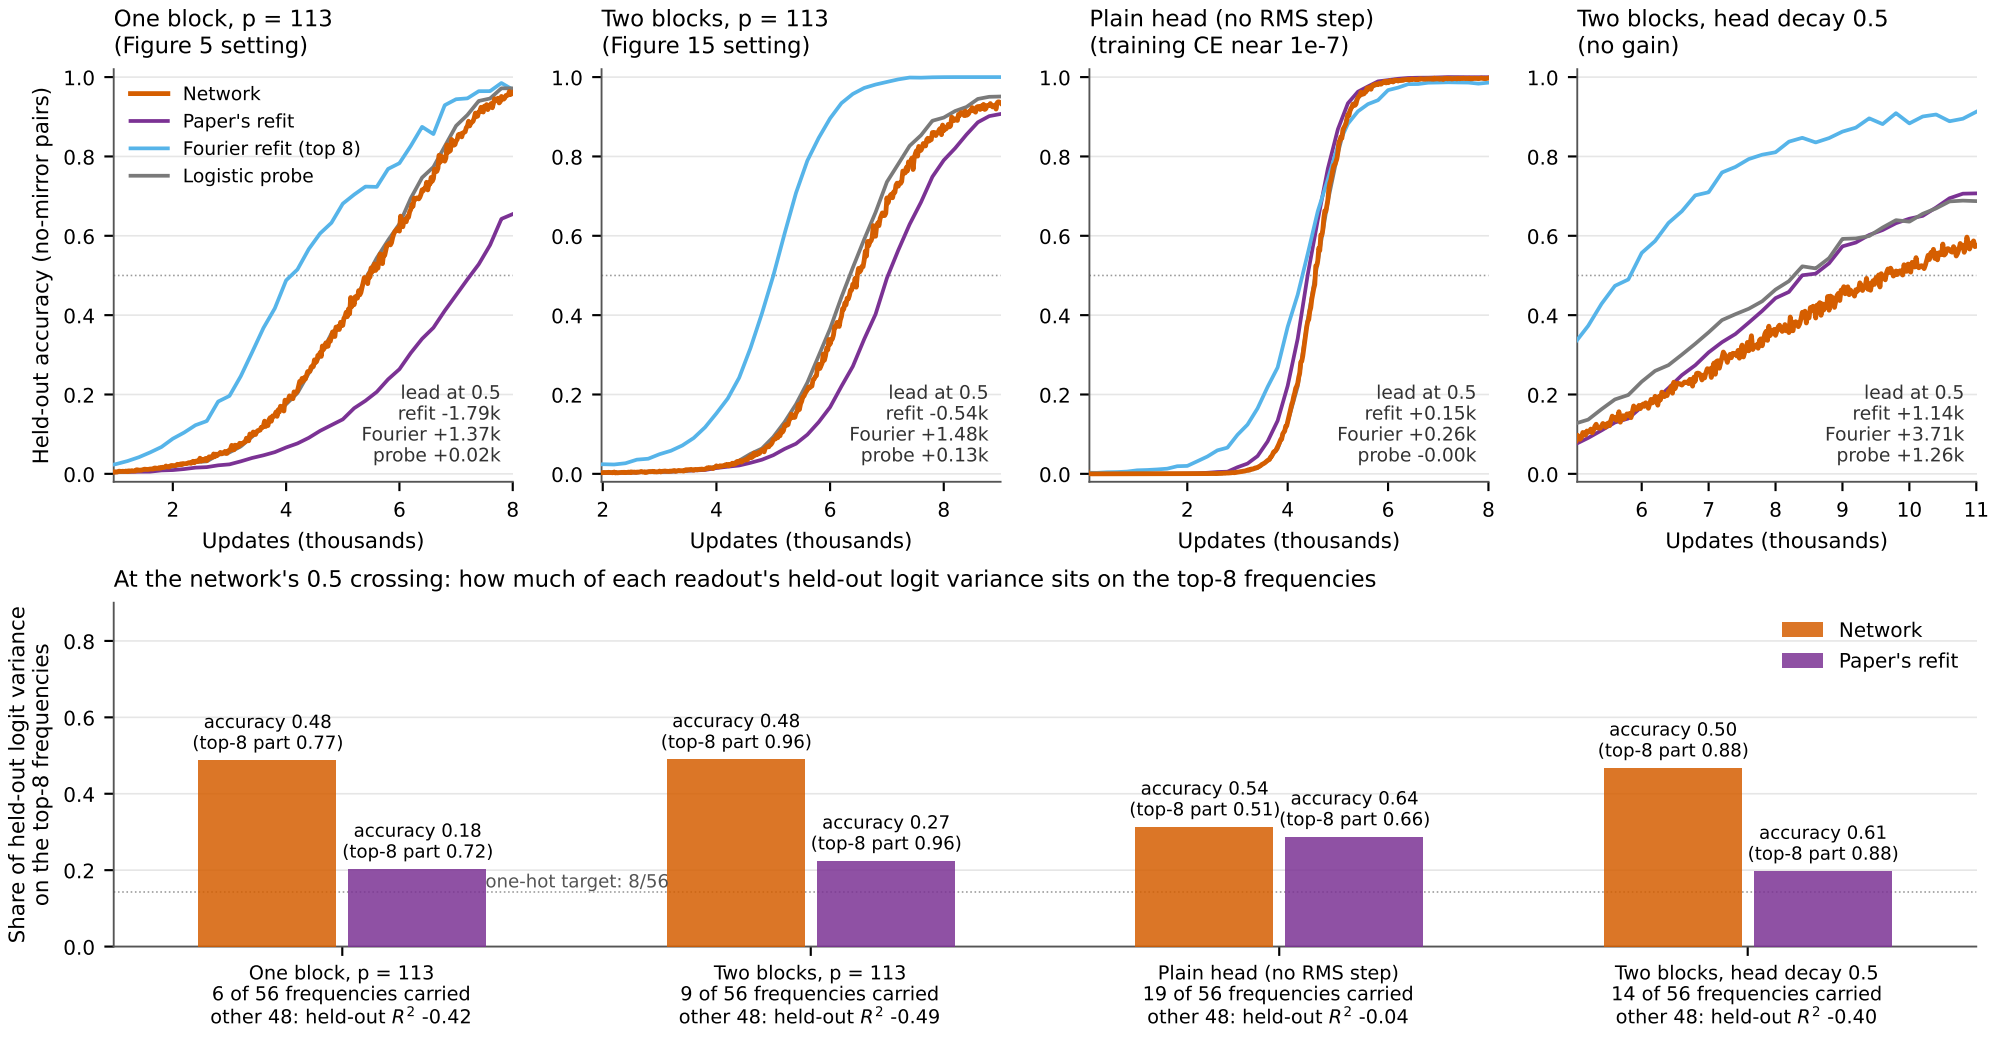
One block, p = 113
(Figure 5 setting)
Two blocks, p = 113
(Figure 15 setting)
Plain head (no RMS step)
(training CE near 1e-7)
Two blocks, head decay 0.5
(no gain)
Network
Paper's refit
Fourier refit (top 8)
Logistic probe
Held-out accuracy (no-mirror pairs)
Updates (thousands)
lead at 0.5
refit -1.79k
Fourier +1.37k
probe +0.02k
lead at 0.5
refit -0.54k
Fourier +1.48k
probe +0.13k
lead at 0.5
refit +0.15k
Fourier +0.26k
probe -0.00k
lead at 0.5
refit +1.14k
Fourier +3.71k
probe +1.26k
At the network's 0.5 crossing: how much of each readout's held-out logit variance sits on the top-8 frequencies
Share of held-out logit variance on the top-8 frequencies
accuracy 0.48
(top-8 part 0.77)
accuracy 0.18
(top-8 part 0.72)
one-hot target: 8/56
accuracy 0.48
(top-8 part 0.96)
accuracy 0.27
(top-8 part 0.96)
accuracy 0.54
(top-8 part 0.51)
accuracy 0.64
(top-8 part 0.66)
accuracy 0.50
(top-8 part 0.88)
accuracy 0.61
(top-8 part 0.88)
One block, p = 113
6 of 56 frequencies carried
other 48: held-out R^2 -0.42
Two blocks, p = 113
9 of 56 frequencies carried
other 48: held-out R^2 -0.49
Plain head (no RMS step)
19 of 56 frequencies carried
other 48: held-out R^2 -0.04
Two blocks, head decay 0.5
14 of 56 frequencies carried
other 48: held-out R^2 -0.40