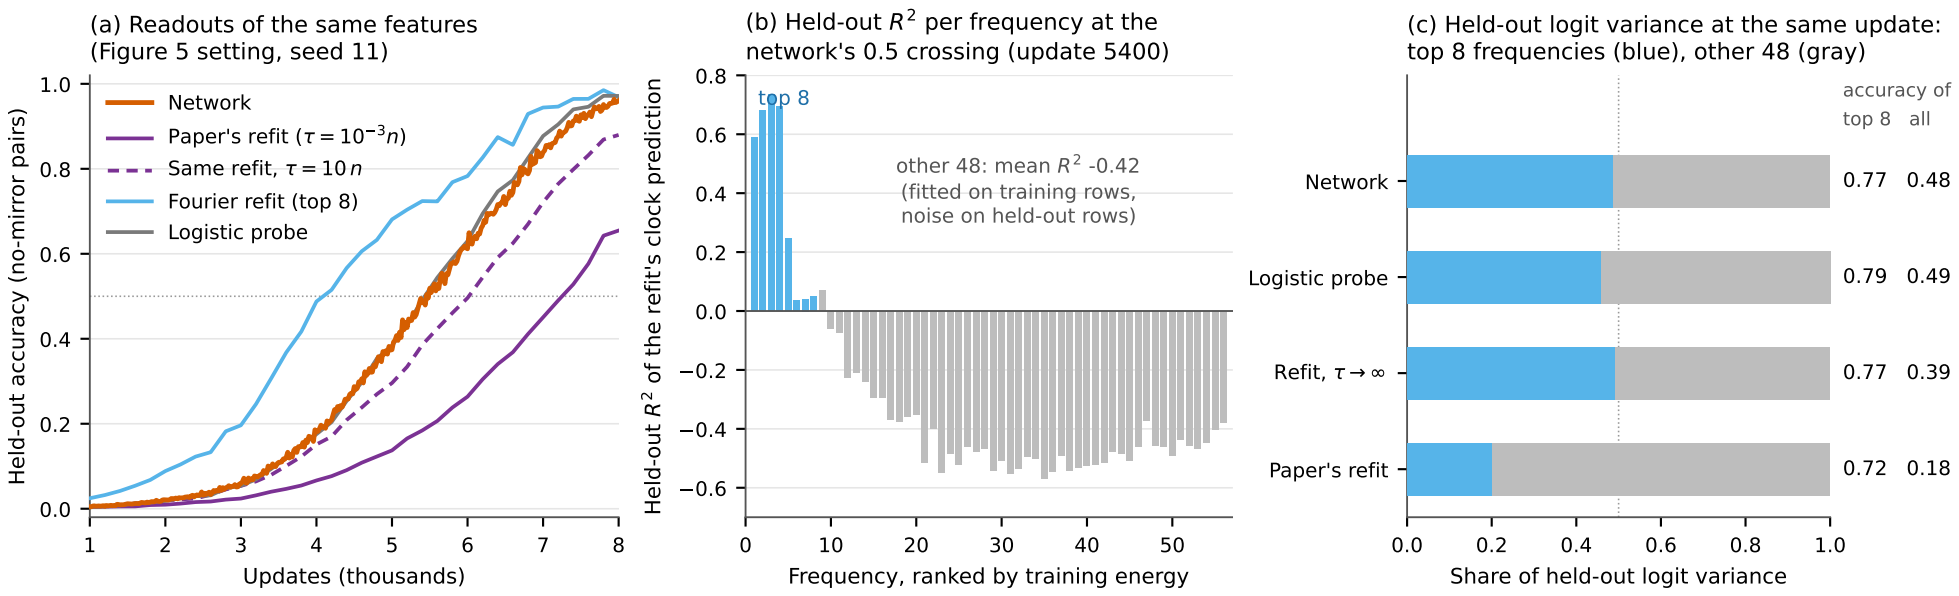
(a) Readouts of the same features
(Figure 5 setting, seed 11)
Network
Paper's refit ($\tau = 10^{-3}n$)
Same refit, $\tau = 10\,n$
Fourier refit (top 8)
Logistic probe
Held-out accuracy (no-mirror pairs)
Updates (thousands)
(b) Held-out R^2 per frequency at the
network's 0.5 crossing (update 5400)
top 8
other 48: mean R^2 -0.42
(fitted on training rows,
noise on held-out rows)
Held-out R^2 of the refit's clock prediction
Frequency, ranked by training energy
(c) Held-out logit variance at the same update:
top 8 frequencies (blue), other 48 (gray)
accuracy of
top 8 all
Network 0.77 0.48
Logistic probe 0.79 0.49
Refit, $\tau \to \infty$ 0.77 0.39
Paper's refit 0.72 0.18
Share of held-out logit variance